

In [1]:

```
1 # Libraries
2 import pandas as pd
3 import os
```

In [2]:

```
1 # importing dataset
2
3 os.chdir('/Users/tomisin/Dropbox/My Mac (Tomisins-MacBook-Pro.local)/Documents/D
```

In [4]:

```
1 # importing dataset
2 df = pd.read_csv('winequality-red.csv')
```

In [5]:

```
1 df
```

Out[5]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcc
0	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	
1	7.8	0.880	0.00	2.6	0.098	25.0	67.0	0.99680	3.20	0.68	
2	7.8	0.760	0.04	2.3	0.092	15.0	54.0	0.99700	3.26	0.65	
3	11.2	0.280	0.56	1.9	0.075	17.0	60.0	0.99800	3.16	0.58	
4	7.4	0.700	0.00	1.9	0.076	11.0	34.0	0.99780	3.51	0.56	
...	
1594	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	0.58	1
1595	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	0.76	1
1596	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	0.75	1
1597	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	0.71	1
1598	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	0.66	1

1599 rows × 12 columns

In [6]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1599 entries, 0 to 1598
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	fixed acidity	1599 non-null	float64
1	volatile acidity	1599 non-null	float64
2	citric acid	1599 non-null	float64
3	residual sugar	1599 non-null	float64
4	chlorides	1599 non-null	float64
5	free sulfur dioxide	1599 non-null	float64
6	total sulfur dioxide	1599 non-null	float64
7	density	1599 non-null	float64
8	pH	1599 non-null	float64
9	sulphates	1599 non-null	float64
10	alcohol	1599 non-null	float64
11	quality	1599 non-null	int64

```
dtypes: float64(11), int64(1)
```

```
memory usage: 150.0 KB
```

In [8]:

```
1 df.isna().sum() #shows columns with their number of missing values
```

Out[8]:

fixed acidity	0
volatile acidity	0
citric acid	0
residual sugar	0
chlorides	0
free sulfur dioxide	0
total sulfur dioxide	0
density	0
pH	0
sulphates	0
alcohol	0
quality	0

```
dtype: int64
```

In [9]:

```
1 df.describe()
```

Out[9]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide
count	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000	1599.000000
mean	8.319637	0.527821	0.270976	2.538806	0.087467	15.874922	46.467000
std	1.741096	0.179060	0.194801	1.409928	0.047065	10.460157	32.895000
min	4.600000	0.120000	0.000000	0.900000	0.012000	1.000000	6.000000
25%	7.100000	0.390000	0.090000	1.900000	0.070000	7.000000	22.000000
50%	7.900000	0.520000	0.260000	2.200000	0.079000	14.000000	38.000000
75%	9.200000	0.640000	0.420000	2.600000	0.090000	21.000000	62.000000
max	15.900000	1.580000	1.000000	15.500000	0.611000	72.000000	289.000000

In [10]:

```
1 df.columns
```

Out[10]:

```
Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',  
      'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',  
      'pH', 'sulphates', 'alcohol', 'quality'],  
      dtype='object')
```


Making plots

In [11]:

```
1 df['fixed acidity'].plot(kind='density')
```

Out[11]:

<AxesSubplot:ylabel='Density'>



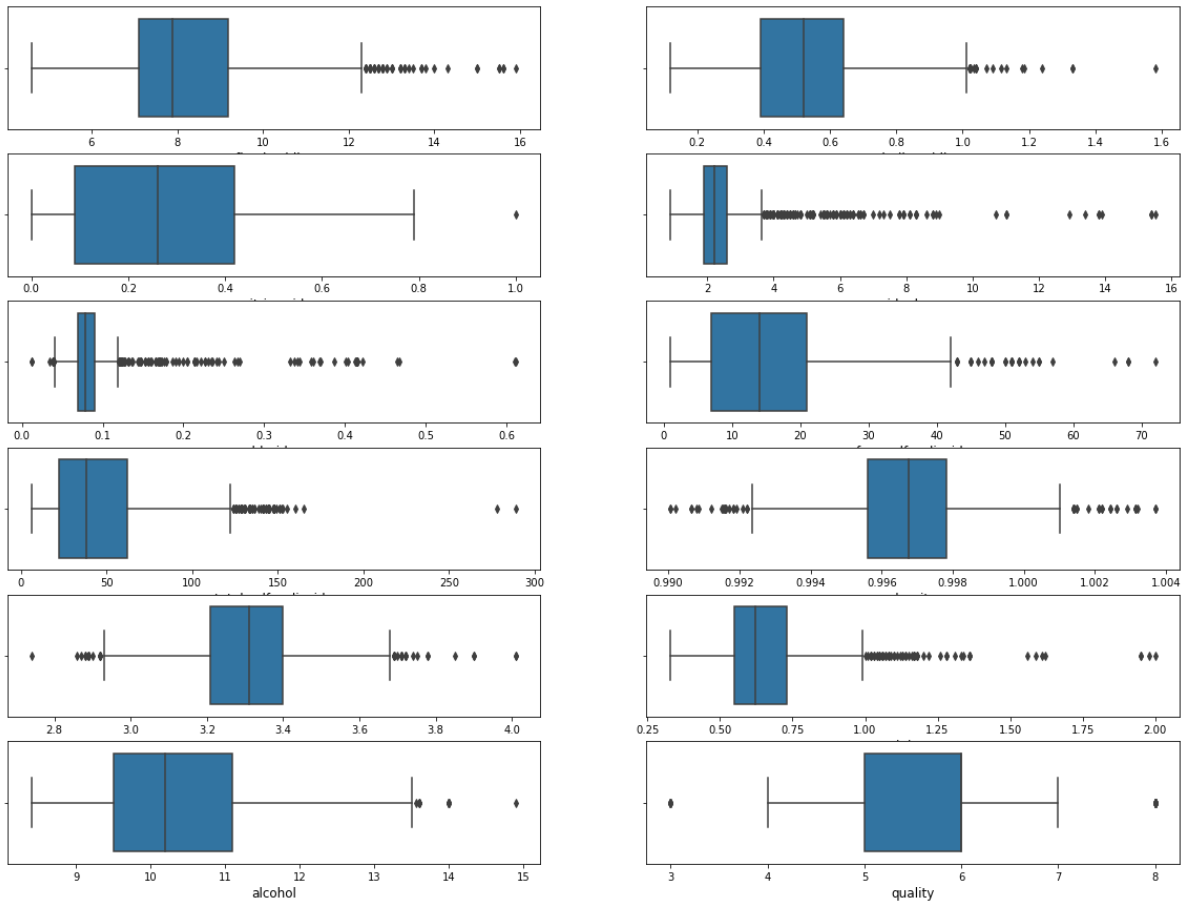
OUTLIERS

In [12]:

```
1 import matplotlib.pyplot as plt
2 import seaborn as sns
3
```

In [13]:

```
1 import matplotlib.pyplot as plt
2 import seaborn as sns
3
4 # Checking for outliers
5 plt.figure(figsize = (20, 15))
6 for i in range (len(df.columns)):
7     plt.subplot(6, 2, i+1)
8     sns.boxplot(x = df.iloc[:, i])
9     plt.xlabel(df.columns[i], size = 12)
```




In [14]:

```
1 df['citric acid'].plot(kind='density')
```

Out[14]:

<AxesSubplot:ylabel='Density'>




In [15]:

```
1 df['density'].plot(kind='density')
```

Out[15]:

<AxesSubplot:ylabel='Density'>



In [16]:

```
1 df.columns[0]
```

Out[16]:

'fixed acidity'

Removing outliers

In [17]:

```
1
2 def Outliers(data, feature):
3     IQ1 = data[feature].quantile(0.25)
4     IQ3 = data[feature].quantile(0.75)
5     IQR = IQ3 - IQ1
6
7     lower_bound = IQ1 - 1.5 * IQR
8     upper_bound = IQ3 + 1.5 * IQR
9
10    index = data.index[ (data[feature] < lower_bound) | (data[feature] > upper_b
11    return index
```

Getting index of all the outliers

In [18]:

```
1
2 index = []
3 for i in df.columns:
4     index.extend(Outliers(df, i))
5 index = set(index)
6 print("Total number of outliers are {}".format(len(index)))
7
8 # Dropping all the outliers
9 df.drop(index, inplace = True, axis = 0)
10 df.shape
```

Total number of outliers are 420

Out[18]:


(1179, 12)

In [19]:

```
1 sns.boxplot(x = df.iloc[:, 2])
```

Out[19]:

<AxesSubplot:xlabel='citric acid'>



Removing outliers

In [20]:

```
1
2 def Outliers(data, feature):
3     IQ1 = data[feature].quantile(0.25)
4     IQ3 = data[feature].quantile(0.75)
5     IQR = IQ3 - IQ1
6
7     lower_bound = IQ1 - 1.5 * IQR
8     upper_bound = IQ3 + 1.5 * IQR
9
10    index = data.index[ (data[feature] < lower_bound) | (data[feature] > upper_b
11    return index
```

In [21]:

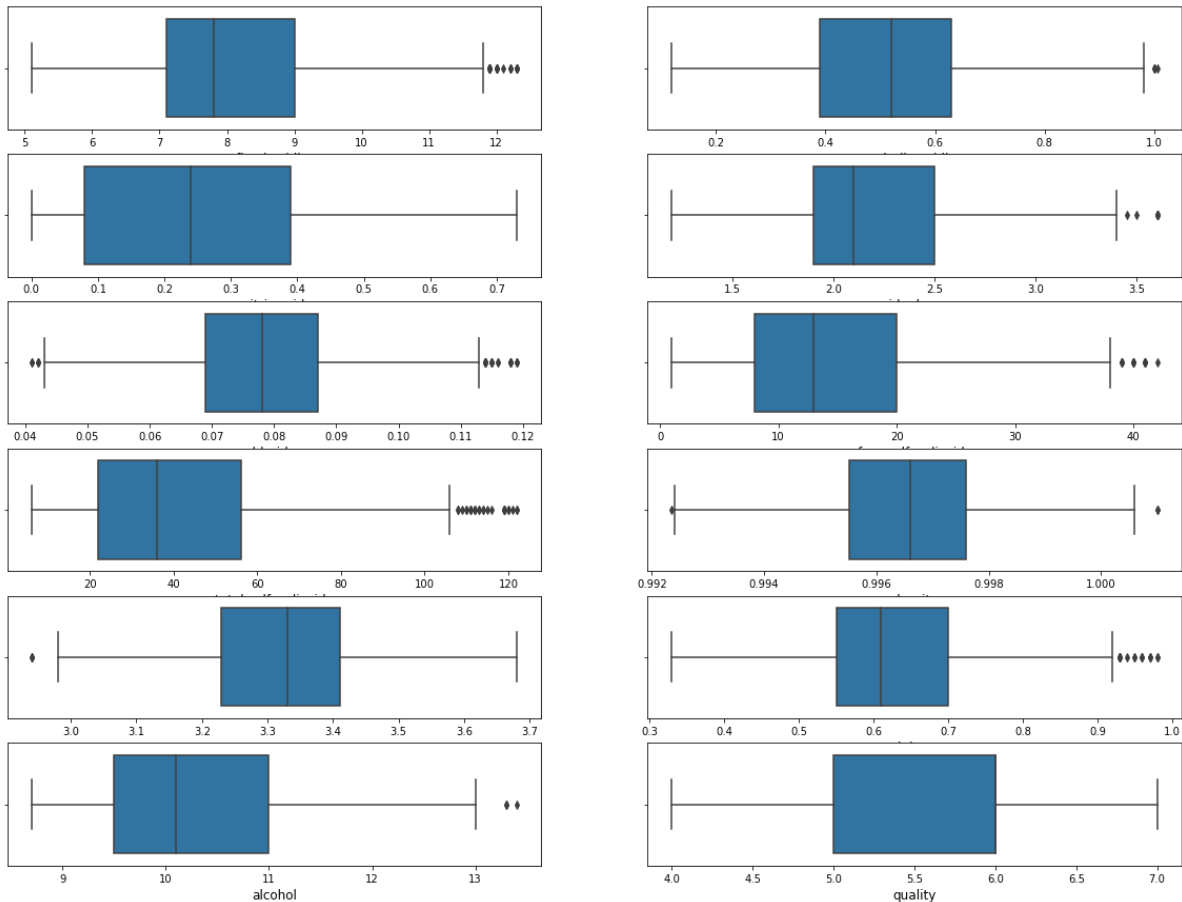
```
1 max(df.quality)
2
```

Out[21]:

7

In [22]:

```
1 import matplotlib.pyplot as plt
2 import seaborn as sns
3
4 # Checking for outliers
5 plt.figure(figsize = (20, 15))
6 for i in range (len(df.columns)):
7     plt.subplot(6, 2, i+1)
8     sns.boxplot(x = df.iloc[:, i])
9     plt.xlabel(df.columns[i], size = 12)
```



Removing outliers and getting index of all the outliers

In [35]:

```
1  # Removing outliers
2  def Outliers(data, feature):
3      IQ1 = data[feature].quantile(0.25)
4      IQ3 = data[feature].quantile(0.75)
5      IQR = IQ3 - IQ1
6
7      lower_bound = IQ1 - 1.5 * IQR
8      upper_bound = IQ3 + 1.5 * IQR
9
10     index = data.index[ (data[feature] < lower_bound) | (data[feature] > upper_b
11     return index
12
13
14
15 # Getting index of all the outliers
16 index = []
17 for i in df.columns:
18     index.extend(Outliers(df, i))
19
20 index = set(index)
21 print("Total number of outliers are {}".format(len(index)))
22
23 # Dropping all the outliers
24 df.drop(index, inplace = True, axis = 0)
25 df.shape
```

Total number of outliers are 2

Out[35]:

(923, 12)

In [25]:

```
1 print(df['quality'].value_counts())
2 plt.figure(figsize = (8, 5))
3 sns.countplot(x = df['quality'])
4 plt.xlabel('Quality', size = 12)
5 plt.ylabel("Count", size = 12)
6 plt.title("Distribution in target variable", size = 12)
```

5 449

6 448


7 124

4 30

Name: quality, dtype: int64

Out[25]:

Text(0.5, 1.0, 'Distribution in target variable')



In [2]:

```
1 !pip install lightgbm
```

```
Requirement already satisfied: lightgbm in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (3.3.3)
Requirement already satisfied: numpy in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (from lightgbm) (1.20.3)
Requirement already satisfied: scikit-learn!=0.22.0 in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (from lightgbm) (0.24.2)
Requirement already satisfied: wheel in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (from lightgbm) (0.37.0)
Requirement already satisfied: scipy in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (from lightgbm) (1.7.1)
Requirement already satisfied: threadpoolctl>=2.0.0 in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (from scikit-learn!=0.22.0->lightgbm) (2.2.0)
Requirement already satisfied: joblib>=0.11 in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (from scikit-learn!=0.22.0->lightgbm) (1.1.0)
```

In [3]:

```
1 !pip install xgboost
```

```
Requirement already satisfied: xgboost in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (1.7.2)
Requirement already satisfied: numpy in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (from xgboost) (1.20.3)
Requirement already satisfied: scipy in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (from xgboost) (1.7.1)
```

In [4]:

```
1 !pip install catboost
```

```
Requirement already satisfied: catboost in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (1.1.1)
Requirement already satisfied: six in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (from catboost) (1.16.0)
Requirement already satisfied: scipy in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (from catboost) (1.7.1)
Requirement already satisfied: matplotlib in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (from catboost) (3.4.3)
Requirement already satisfied: graphviz in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (from catboost) (0.20.1)
Requirement already satisfied: pandas>=0.24.0 in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (from catboost) (1.3.4)
Requirement already satisfied: numpy>=1.16.0 in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (from catboost) (1.20.3)
Requirement already satisfied: plotly in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (from catboost) (5.11.0)
Requirement already satisfied: python-dateutil>=2.7.3 in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (from pandas>=0.24.0->catboost) (2.8.2)
Requirement already satisfied: pytz>=2017.3 in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (from pandas>=0.24.0->catboost) (2021.3)
Requirement already satisfied: cyclor>=0.10 in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (from matplotlib->catboost) (0.10.0)
Requirement already satisfied: pyparsing>=2.2.1 in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (from matplotlib->catboost) (3.0.4)
Requirement already satisfied: kiwisolver>=1.0.1 in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (from matplotlib->catboost) (1.3.1)
Requirement already satisfied: pillow>=6.2.0 in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (from matplotlib->catboost) (8.4.0)
Requirement already satisfied: tenacity>=6.2.0 in /Users/tomisin/opt/anaconda3/lib/python3.9/site-packages (from plotly->catboost) (8.1.0)
```

In [8]:

```
1 import matplotlib.pyplot as plt
2 import seaborn as sns
3 import warnings
4 warnings.filterwarnings('ignore')
5 from sklearn.ensemble import RandomForestClassifier
6 from sklearn.tree import DecisionTreeClassifier
7 from sklearn.linear_model import LogisticRegression
8 from sklearn.naive_bayes import GaussianNB
9 from sklearn.neighbors import KNeighborsClassifier
10 from sklearn.svm import SVC
11 from sklearn.model_selection import train_test_split
12 from sklearn.preprocessing import LabelEncoder, StandardScaler
13 from sklearn.metrics import accuracy_score, auc, roc_curve, roc_auc_score, mean
```

In [41]:

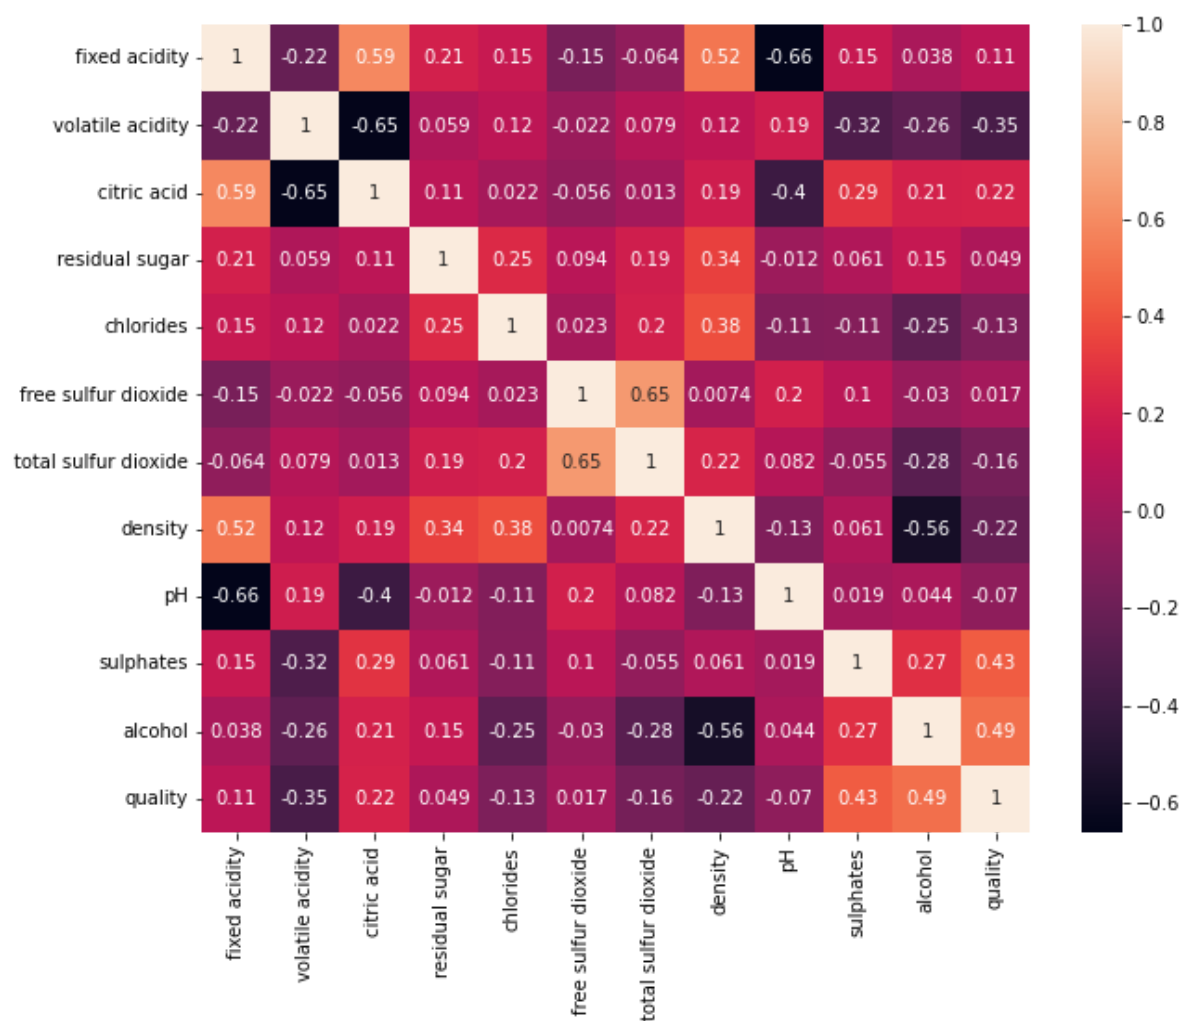
```
1 data = df
```

In [42]:

```
1 cor = data.corr()
2 plt.figure(figsize = (10,8))
3 sns.heatmap(cor, annot = True)
```

Out[42]:

<AxesSubplot:>



In [43]:

```
1 x = data.iloc[:, :-1]
2 x.head()
```

Out[43]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4
1	7.8	0.88	0.00	2.6	0.098	25.0	67.0	0.9968	3.20	0.68	9.8
2	7.8	0.76	0.04	2.3	0.092	15.0	54.0	0.9970	3.26	0.65	9.8
3	11.2	0.28	0.56	1.9	0.075	17.0	60.0	0.9980	3.16	0.58	9.8
4	7.4	0.70	0.00	1.9	0.076	11.0	34.0	0.9978	3.51	0.56	9.4

In [44]:

```
1 y = data.iloc[:, -1]
2 y.head()
```

Out[44]:

```
0    5
1    5
2    5
3    6
4    5
Name: quality, dtype: int64
```

In []:

```
1
```

In []:

```
1
```

In []:

```
1
```

In []:

```
1
```

In []:

```
1
```